



The three-minute pitch

What to say, and what to have on screen while you say it.

Seven beats, 180 seconds, about 400 spoken words. The words are written to be said rather than read - short sentences, the number before the explanation, and nothing you would not say out loud. Every figure in bold comes from the repository's own benchmark output, so this script cannot quote a number the project no longer produces.

Rs 1.23 cr

recovered per 2 days, mean of 8 seeds

38.2%

of the money the incidents put at risk

0/8

seeds where the router lost money

Before you press record

Two minutes of setup that will save you a retake.

- **Open the console first and let it settle.** `localhost:8501`. The default seed is already warmed into the image, so it opens in about a second - but open it before you start recording anyway, so the first thing on camera is the board and not a spinner.
- **Browser at 100% zoom, window maximised.** The header is built to fill the width; at 80% the tagline and the quotes crowd each other.
- **Have three tabs ready** so you never wait for a page: the console on Overview, a second on Decision record, and a terminal.
- **Do not use --fast on the showcase** if you show it. It runs one simulated day and its figures are smaller than the ones in this script, which are two-day numbers.
- **Say the eight-seed figure, not the single run.** The tile says Rs 1.19 cr for this run; the sidebar says **Rs 1.23 cr** across 8 seeds with a confidence interval. The second one is the defensible number, so it is the one in the script.

The shape of the three minutes

Time	Beat	On screen
0:00	The problem nobody gets paged for	Board / masthead
0:25	Watch it catch one	Decision record, top card
0:53	What it did, and what that was worth	Same card, then tiles
1:21	The half a dashboard never shows	Refused, then Rollbacks
1:49	Where I was wrong, and the measurement	Settings, one click
2:21	One run proves nothing	Sidebar, then How it was measured
2:45	Where the AI is, and where it is not	How it was measured

The script

Bold is a number — land on it, then pause. Square brackets are stage directions, not words. The whole thing is **400 words**, which is about 160 seconds at a normal pace; the rest of the three minutes is the pauses and the clicking.

0:00 The problem nobody gets paged for · 25 sec

ON SCREEN The console open on Overview. Do not scroll yet.

A payment gateway is not one pipe. Traffic splits across acquiring banks, methods and issuing banks — and when payments fail, it is almost never everywhere at once. It is narrow. Say one slice is four percent of volume and it collapses from ninety-six percent to forty. Your headline success rate moves **two points**. Nobody is paged, and the money leaves quietly.

0:25 Watch it catch one · 28 sec

ON SCREEN Click DECISION RECORD. Stop on the top card.

This is RazorGuard. Two simulated days, **1.97 million** payment attempts, **eighteen** injected failures. Here is the worst. A gateway outage: those routes fell from **ninety-four percent to thirty-three**, caught in **one minute**, with **65 lakh** at risk. [point at the blue cells] Everything blue was measured during the run. Everything else is policy written before it. Different kinds of claim, so they do not look alike.

0:53 What it did, and what that was worth · 28 sec

ON SCREEN Scroll to 'What the system did', then back up to the tiles.

It moved traffic off the failing routes in **twenty-two** bounded steps, always leaving a **three percent** canary so the broken route stays observable. Across the run that is **Rs 1.19 cr** recovered, net of fees. And that is a difference between two runs that saw identical demand — not a projection. If the two arms ever diverge, it refuses to report a figure at all.

1:21 The half a dashboard never shows · 28 sec

ON SCREEN Click REFUSED. Then click ROLLBACKS.

This is the part I would want to be asked about. **Ninety-three** times it declined to move money, each with the rule that stopped it. [click Rollbacks] **Thirty-two** times it moved traffic, measured the destination, found it worse, and put the traffic back. Issuer-wide faults it escalates instead — when the bank is failing everywhere, shuffling between broken routes recovers nothing.

1 : 49 **Where I was wrong, and how I know** · 32 sec

ON SCREEN Click SETTINGS → 'Things worth trying' → **Max shift 40%**. Let it re-run on camera.

Everything here is tunable and it really re-runs. The cap on how much one action may move started at forty percent, because forty felt careful. [the page returns] Recovery falls from **1.20 crore to 71 lakh**. My caution was costing more than half the recovery, and only measurement found that. A second sweep found worse: on a saturated fleet, rerouting *loses* money. So there is a breaker that halts the system when its own shifts stop paying.

Click the link just before you start talking, so the re-run happens under the first sentence rather than in silence.

2 : 21 **One run proves nothing** · 24 sec

ON SCREEN Point at the sidebar 'Validated result' block.

One run proves nothing, so — **8 seeds**. Mean **Rs 1.23 cr**, ninety-five percent interval **Rs 1.18 cr to Rs 1.28 cr**. That is **38.2%** of the money those incidents put at risk, and **0.58 points** of fleet success rate. **Zero of 8** lost money. And the detectors cannot import the file that defines the incidents — a test fails if anyone tries.

2 : 45 **Where the AI is, and where it is not** · 14 sec

ON SCREEN Click HOW IT WAS MEASURED. Land on the bottom row.

Detection, attribution, policy and routing are deterministic. A model writes the operator note and explains past decisions — and can change none of them. An LLM that can move money is a liability. Thank you.

If you are running long

Cut in this order. Each of these can go without breaking the argument, because the beat before it already made the point.

Cut	Saves	What you lose
The rollbacks half of beat 4 (1:21)	8 sec	One example of self-correction. The refusals still carry it.
The capacity sweep sentence in beat 5 (1:49)	9 sec	The second finding. The shift-cap number is the stronger one.
The confidence interval in beat 6 (2:21)	6 sec	Say the mean and 'zero of eight lost money' instead.

If you are running short

Open the incident replay and read two lines of the ledger aloud — a refusal and the rollback that follows it. It is the most convincing twenty seconds in the project, and it is real output rather than anything composed for a demo.

Questions you should expect

Short answers, because a long one sounds rehearsed.

Question	Answer
Razorpay already has Optimizer. Why this?	Optimizer routes. This measures whether the routing paid, refuses when it would not, and writes down why. The evidence half is the product.
It is a simulator. Why should I believe the number?	Because the number is a difference between two runs that saw identical demand, the detectors cannot see the incident plan, and the run refuses to report at all if the arms diverge. Three tests enforce those.
Where is the AI?	Deliberately out of the loop that moves money. It writes the operator note and explains a past decision. Detection, attribution and policy are deterministic and auditable.
Does it work on real traffic?	The service takes real payment outcomes at POST /ingest, authenticated. The executor has run live against a Razorpay test account — six of six orders created, with the decision context in the notes.
What breaks first at scale?	Attribution needs a global view, so it is one active instance with lease-based failover rather than shards. That is in DEPLOY.md, with what is still open.

The numbers, in one place

Everything below is interpolated from `bench/results/*.json` when this document is built, so it matches what the console will show on camera.

Figure	Value	Where it comes from
Recovered, mean of 8 seeds	Rs 1.23 cr	validate.py, 8 seeds x 2 days
95% confidence interval	Rs 1.18 cr to Rs 1.28 cr	the same run
Share of money at risk recovered	38.2%	control-arm shortfall as denominator
Success-rate gain	+0.58pp	treatment minus control
Seeds that lost money	0 of 8	validate.py

This run, net of fees	Rs 1.19 cr	experiment.py, seed 7
This run, payments saved	11,452	the same run
Payment attempts, per arm	1,966,458	identical in both arms, or nothing is reported